

Bridging the Archipelago between Row-Stores and Column-Stores for Hybrid Workloads



J Arulraj, et al.

SIGMOD '20


Speaker:

Introduction


Physical Memory layout



Physical Memory layout




?




?

Physical Memory layout



Row Storage


tuple-centric




Column Storage

Columnar

Physical Memory layout




Row Storage



Column Storage


When?

Physical Memory layout



Row Storage


When? Insert, Update




Column Storage

Search for an attribute

Physical Memory layout



Row Storage



Column Storage

When?


Insert, Update

Hotly updated

Search for an attribute


Coldly stored

Physical Memory layout



Row Storage


When? Insert, Update




Column Storage

Search for an attribute

Physical Memory layout



Row Storage



Column Storage

When?


Insert, Update

Transactional

Search for an attribute

Analytical

Physical Memory layout




Row Storage

When?

Insert, Update

OLTP




Column Storage

Search for an attribute

OLAP

Physical Memory layout




Row Storage

When?


Insert, Update

OLTP




Search for an attribute

Physical Memory layout



Row Storage




When?

Insert, Update

Search for (a1,a2) attribute

OLTP

Physical Memory layout




Row Storage

When?

Insert, Update


OLTP



Search for (a1,a2) attribute

HTAP

Physical Memory layout




Row Storage

When?

Insert, Update

OLTP




??

Search for (a1,a2) attribute


HTAP

Physical Memory layout

OLTP




OLAP



Physical Memory layout

OLTP → HTAP ← OLAP



Physical Memory layout



OLTP → HTAP ← OLAP

Physical Memory layout

HTAP


Physical Memory layout

HTAP



Physical Memory layout

HTAP



Which Memory layout

Which?

Which Memory layout

Which?

Do we have any
observation?

Which Memory layout

Which?

Do we have any
observation?

Which Memory layout

Hot data → row

Cold data → column

Which Memory layout

Record Query types?

Which Memory layout Record Query types?

If we know what types of queries there are, then
we definitely can make a good design.

Which Memory layout Record Query types?

If we know what types of queries there are, then
we definitely can make a good design.

But ...

Which Memory layout Record Query types?

If we know what types of queries there are, then
we definitely can make a good design.

But ...

Can we be smarter?

Which Memory layout

Record Query types?

Self-adaptive?

Self-adaptive algorithm



Record

Self-adaptive algorithm




Record




Analysis

Self-adaptive algorithm




Self-adaptive algorithm




On-line k-means algorithm

Find
Representatives



On-line k-means algorithm

Find
Representatives



1st step:


Using recent n Queries,
Find k Representatives R

2nd step:


Generate vertical partitioned
layout using R with greedy
algorithm. (Largest cluster first)

On-line k-means algorithm

Find
Representatives




FSM (Flexible Storage Model)




Different tables, different vertical layouts

On-line k-means algorithm

Find
Representatives




FSM (Flexible Storage Model)




Different layouts → different access methods?

On-line k-means algorithm

Find
Representatives



FSM (Flexible Storage Model)




Different layouts → different access methods?

Inefficient !!




On-line k-means algorithm

Find
Representatives




FSM (Flexible Storage Model)




Provide Abstract layer !!

On-line k-means algorithm

Find
Representatives



FSM (Flexible Storage Model)



Provide Abstract layer !!

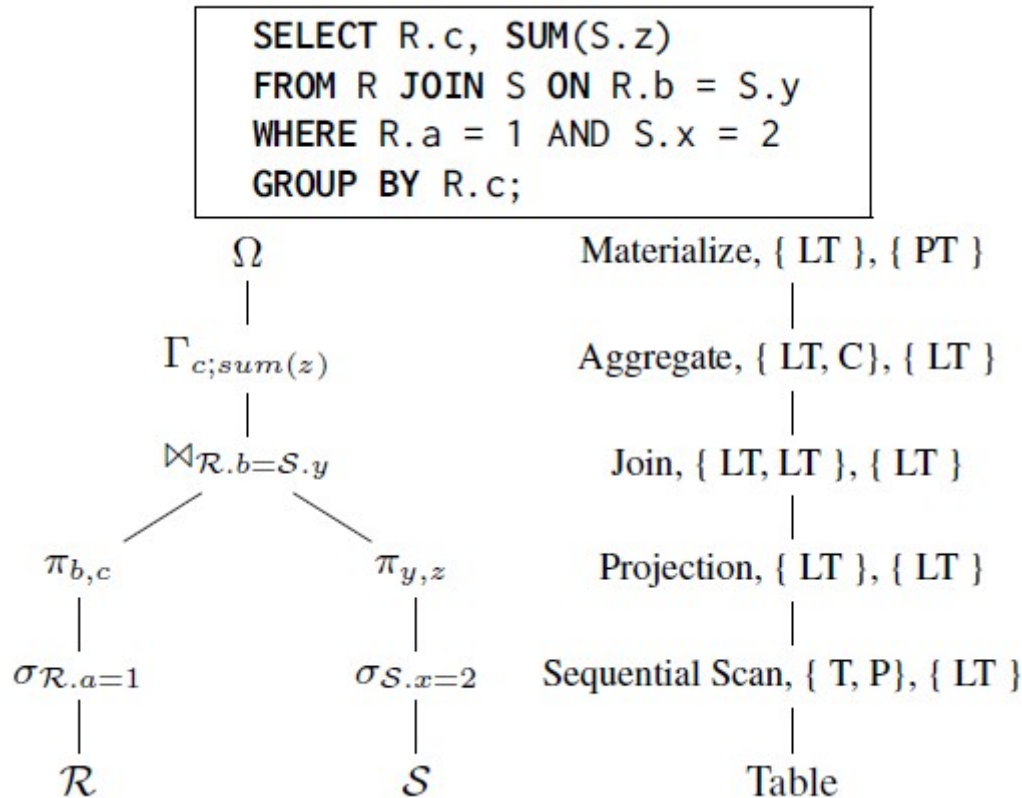


One mutual access interface.

Abstraction

```
SELECT R.c, SUM(S.z)
FROM R JOIN S ON R.b = S.y
WHERE R.a = 1 AND S.x = 2
GROUP BY R.c;
```

Abstraction



LT : logical tile

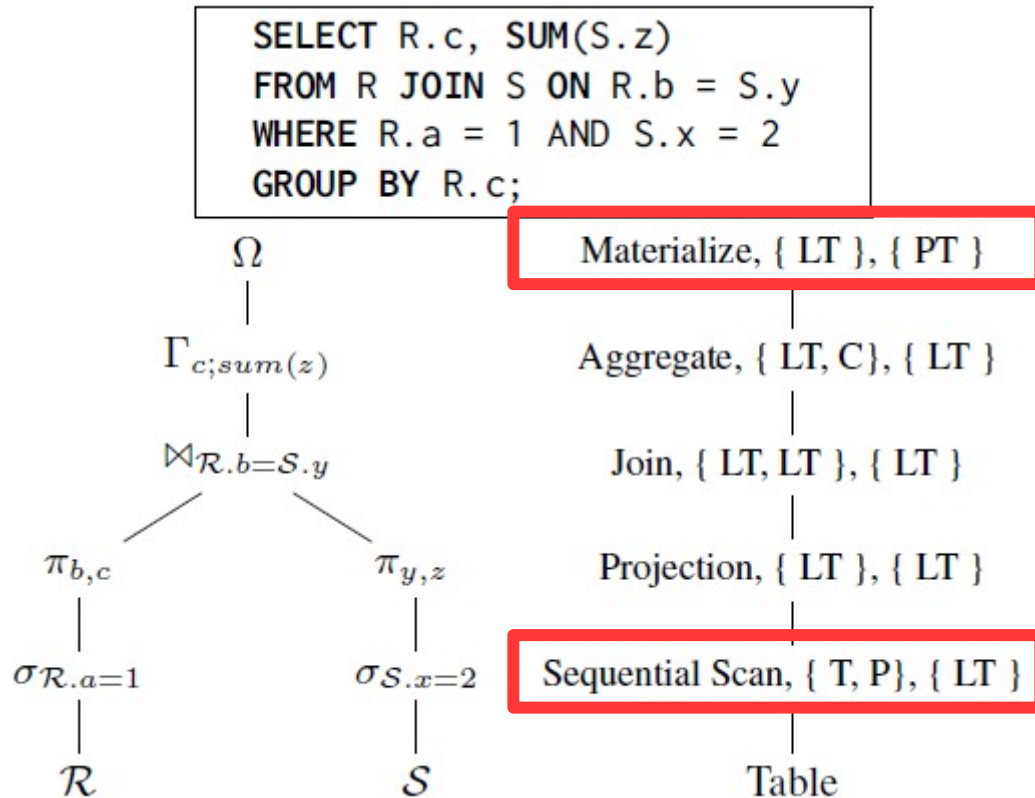
PT : physical tile

T : table

Attributes : C

Predicate : P

Abstraction



LT : logical tile

PT : physical tile

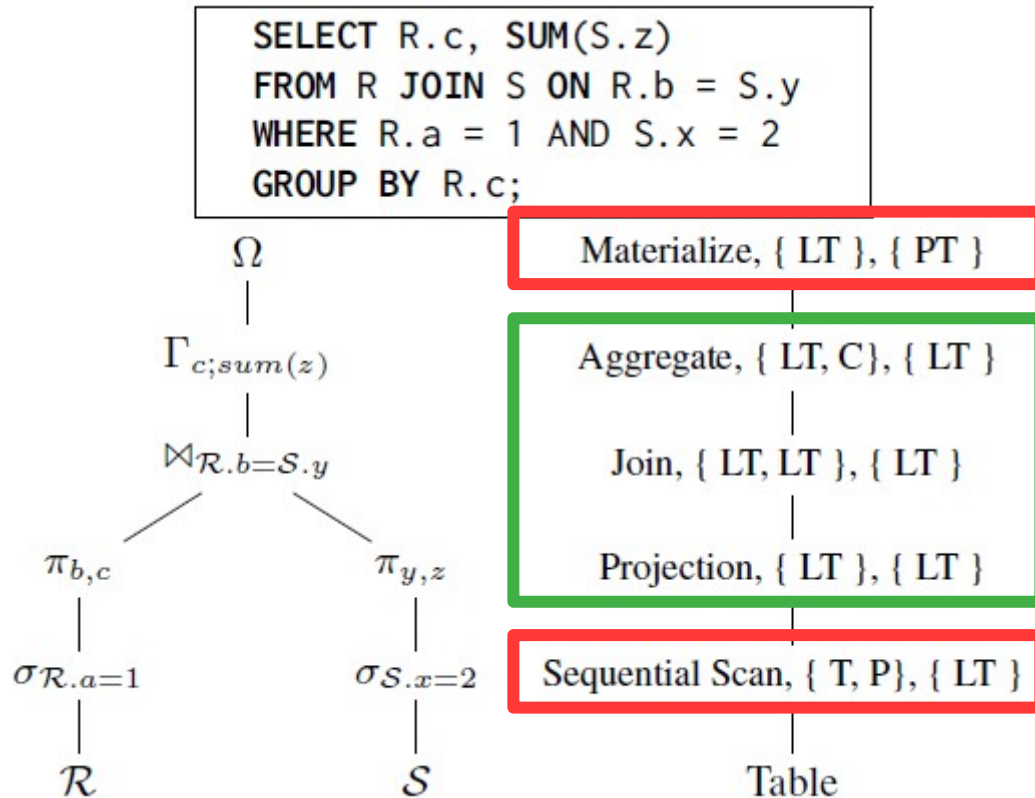
T : table

Attributes : C

Predicate : P

Physical data

Abstraction



LT : logic tile

PT : physical tile

T : table

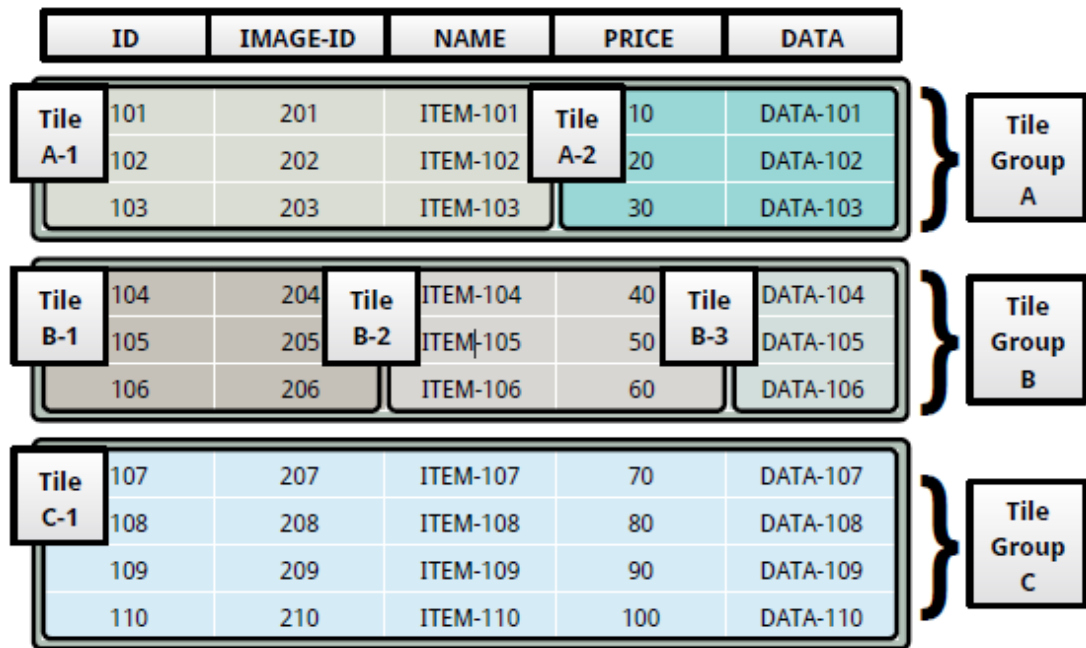
Attributes : C

Predicate : P

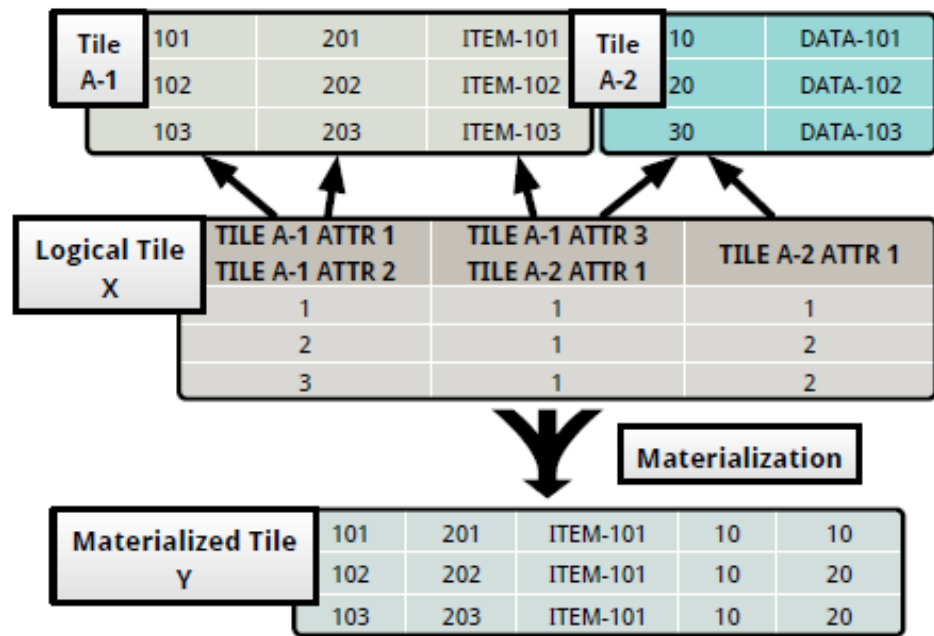
logical data +

Flexible materialization

Tiles



Physical Tile



Logical Tile

Parallelism

Definitely we do not want stalling.

Parallelism

Definitely we do not want stalling.

Write while reconfigure memory layout

Parallelism

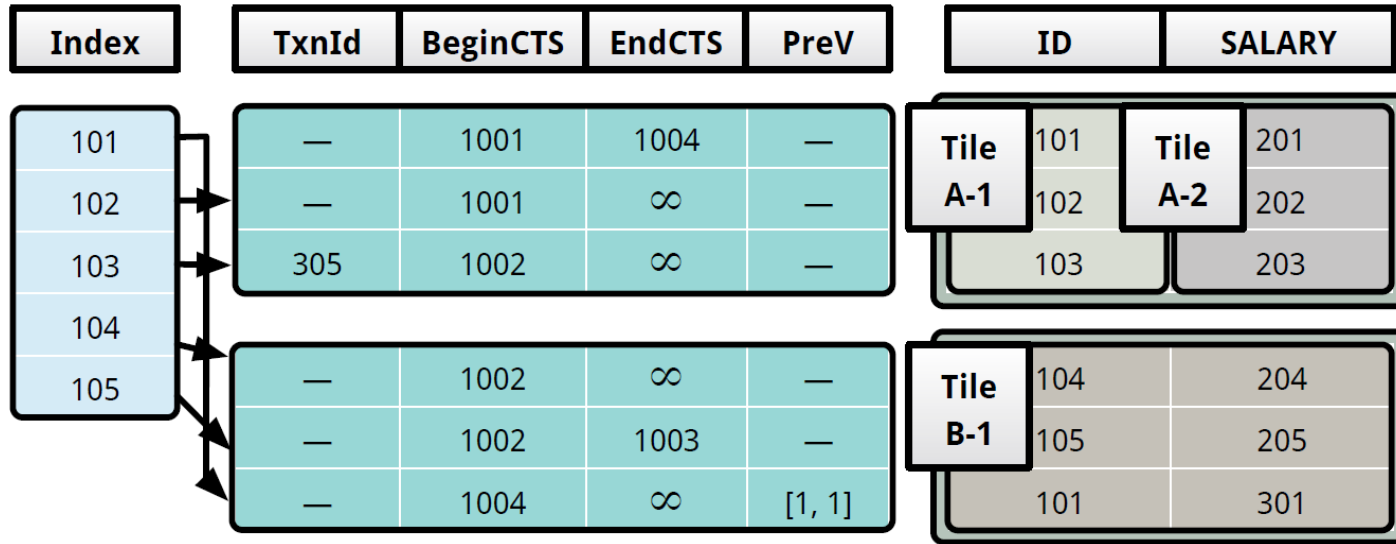
Definitely we do not want stalling.

Write while reconfigure memory layout

→ MVCC

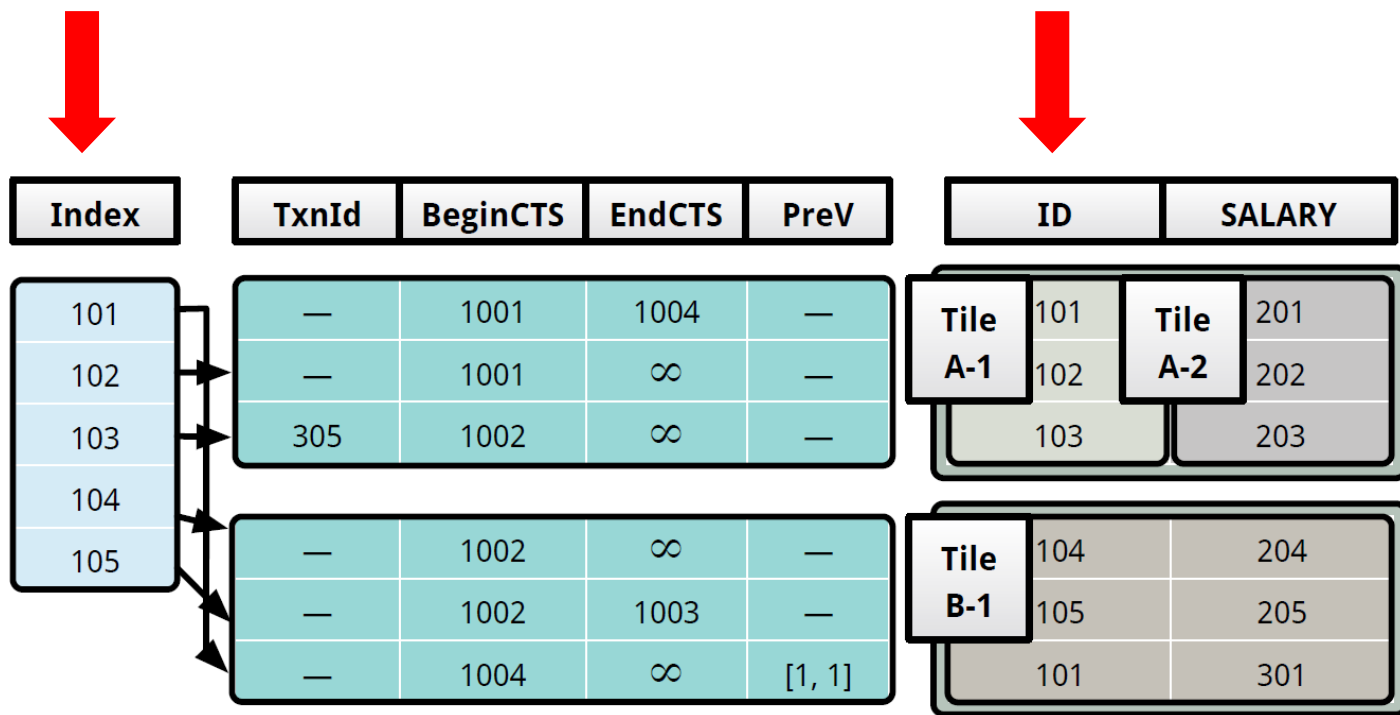
(Multi-Version Concurrency Control)

MVCC



Metadata for MVCC

MVCC



Metadata for MVCC

MVCC

Unique transaction identifier




Running
transaction

Index	TxnId	BeginCTS	EndCTS	PreV	ID	SALARY
101	—	1001	1004	—	Tile A-1 101 102 103	Tile A-2 201 202 203
102	—	1001	∞	—		
103	305	1002	∞	—		
104	—	1002	∞	—	Tile B-1 104 105 101	204 205 301
105	—	1002	1003	—		
	—	1004	∞	[1, 1]		

Metadata for MVCC

MVCC

Unique commit timestamp



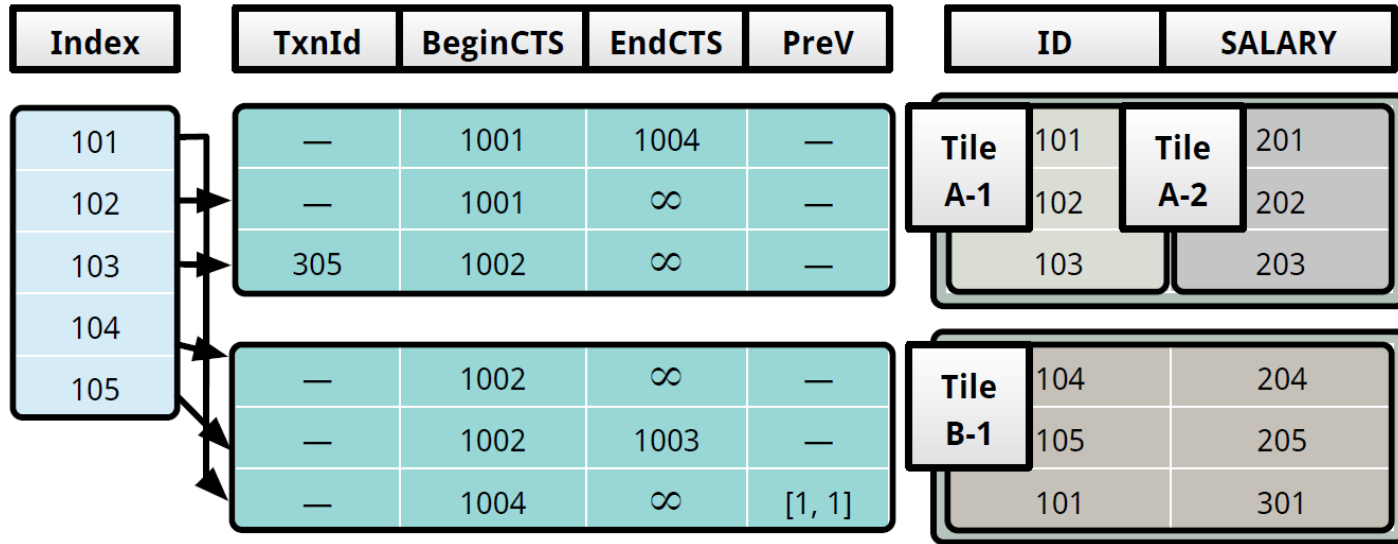
Metadata for MVCC

MVCC

i) Update



(atomic)




Metadata for MVCC

MVCC

i) Update

↓ (after commit)



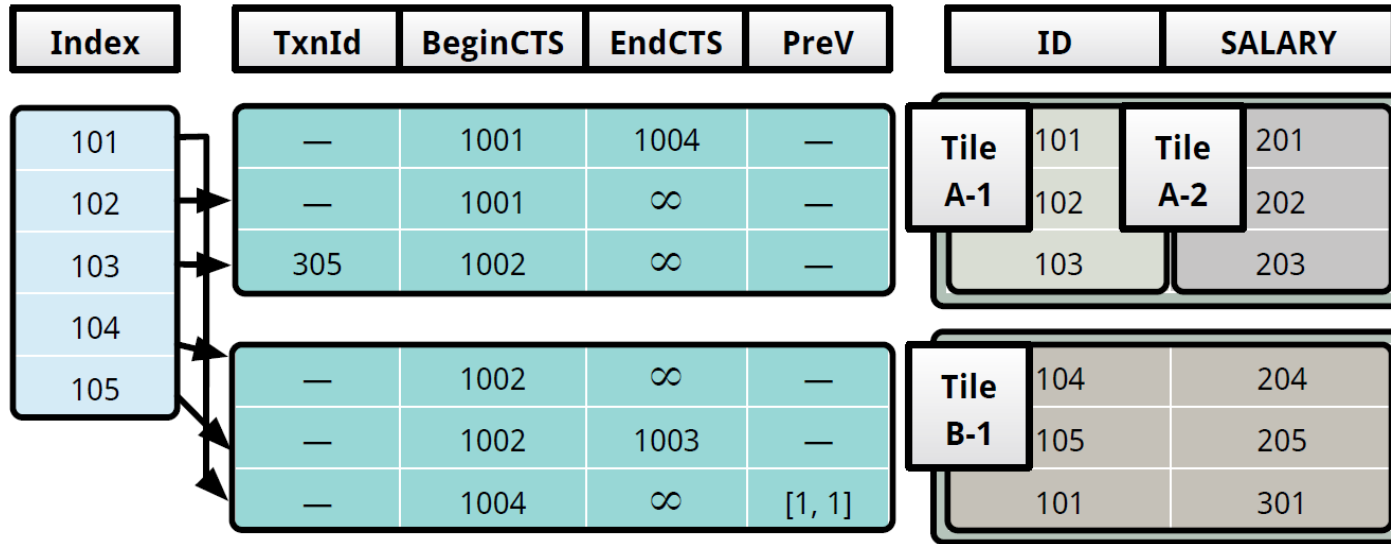
Metadata for MVCC

MVCC

ii) Delete



(atomic)




Metadata for MVCC

MVCC



(after commit)

ii) Delete

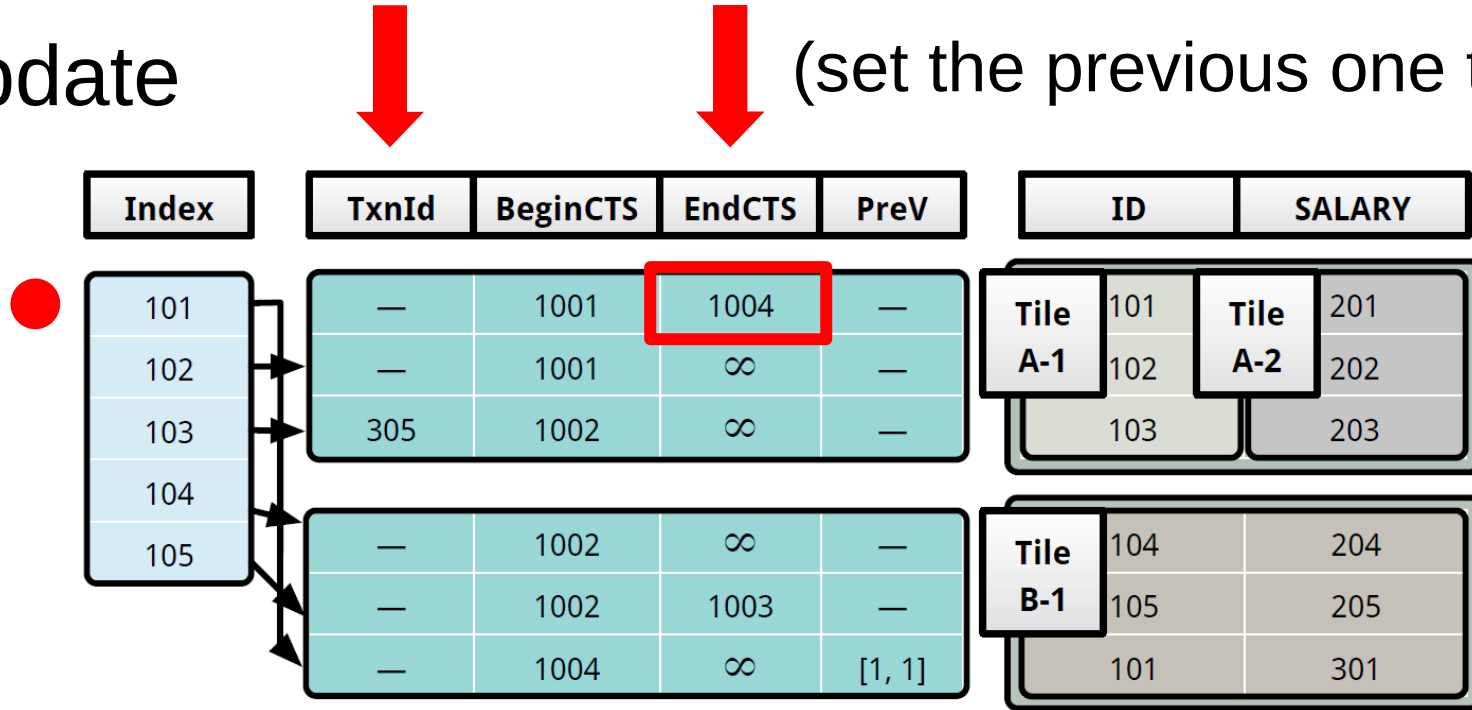


Metadata for MVCC

MVCC

iii) Update

(set the previous one to invalid)

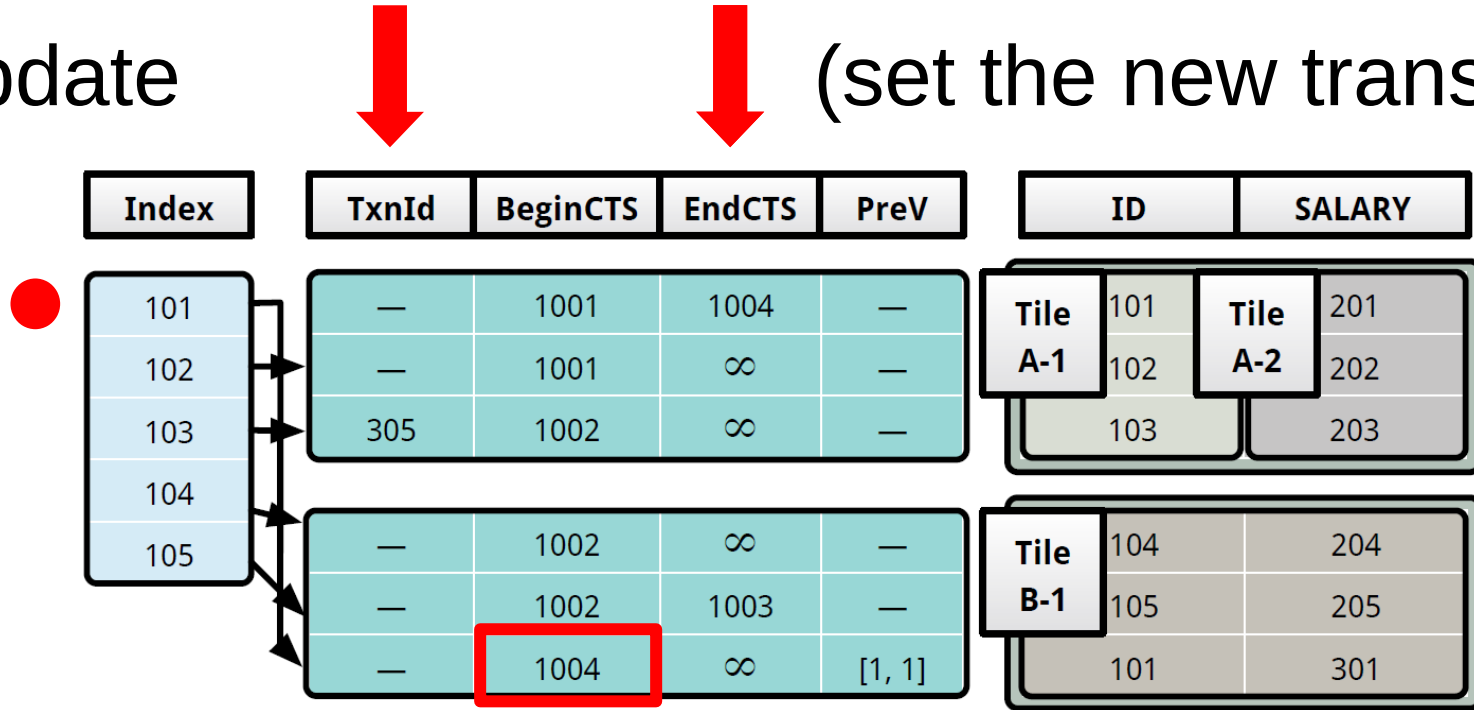


Metadata for MVCC

MVCC

iii) Update

(set the new transaction)

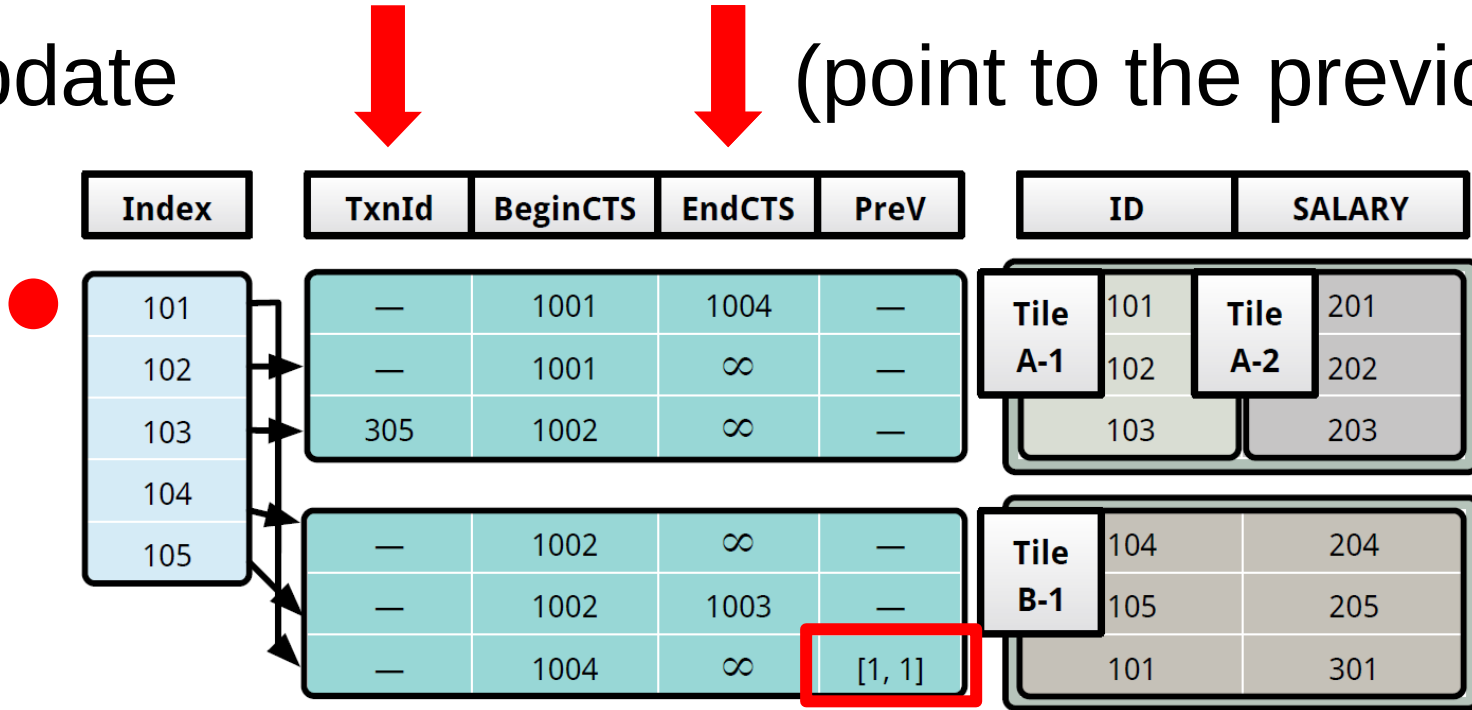


Metadata for MVCC

MVCC

iii) Update

(point to the previous one)




Metadata for MVCC

MVCC

iii) Update

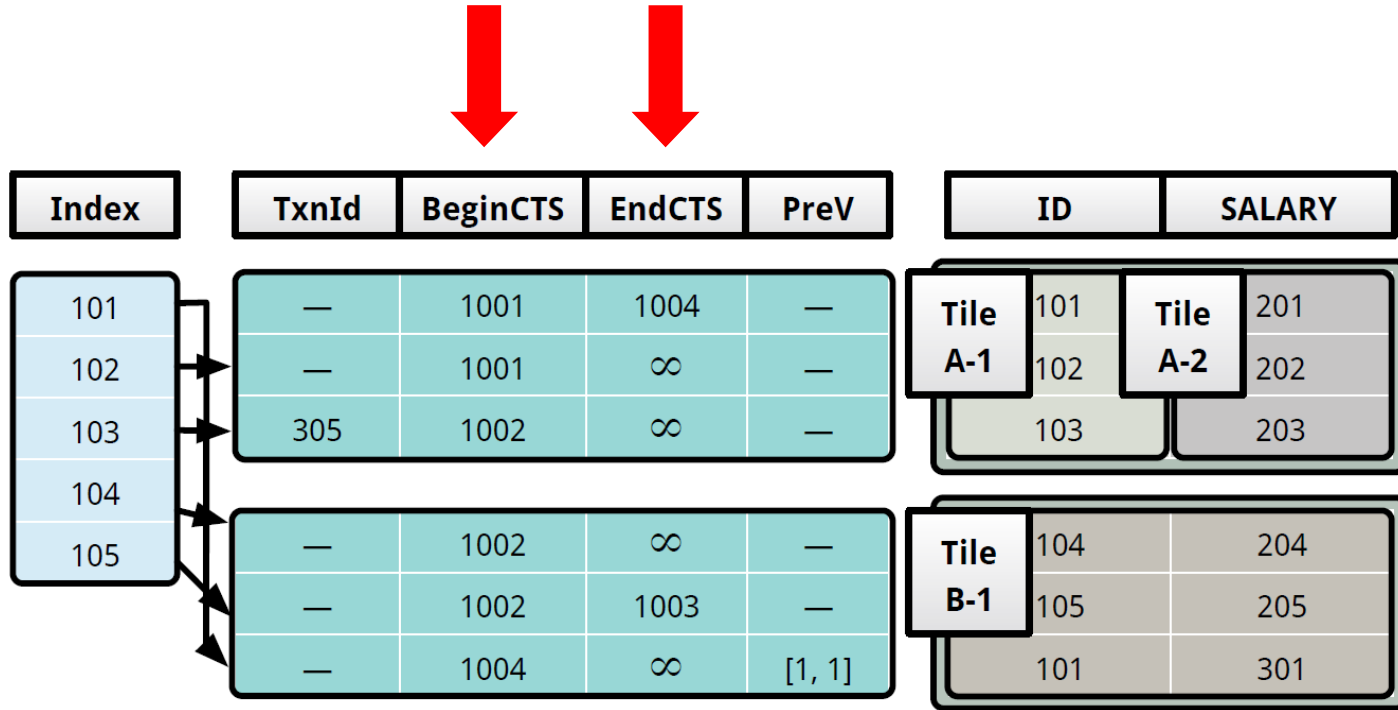
(point to the previous one)



Metadata for MVCC

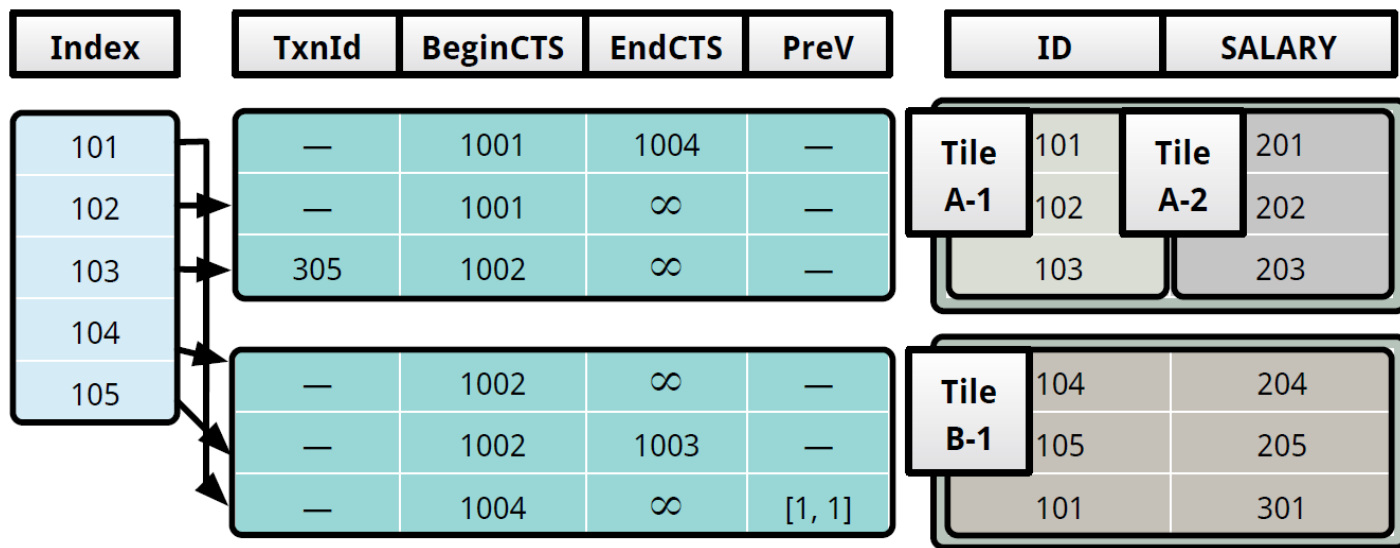
MVCC

Search



Looking through the metadata, find the one within its visibility ($\text{BeginCTS} < \text{transaction ID} < \text{EndCTS}$)

MVCC



Garbage Collector claims those old spaces back

Performance

What about the overhead?

Is this method practical?

Performance


What about the overhead?

Is this method practical?


Let us see the results.

Evaluation


NSM / DSM / FSM SCAN




(a) Scan, Narrow, Read Only



(b) Scan, Narrow, Hybrid




(c) Scan, Wide, Read Only




(d) Scan, Wide, Hybrid

Storage Models : NSM DSM FSM


NSM / DSM / FSM Aggregate




(e) Aggregate, Narrow, Read Only



(f) Aggregate, Narrow, Hybrid




(g) Aggregate, Wide, Read Only




(h) Aggregate, Wide, Hybrid

Storage Models : NSM DSM FSM


Adaptation




Logical tiles




(a) Scan, Narrow, Read Only




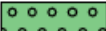

(b) Scan, Narrow, Hybrid



(c) Scan, Wide, Read Only



(d) Scan, Wide, Hybrid

Storage Models :  NSM  DSM  FSM

Conclusion

Conclusion

- Memory Layout (column storage, row storage)
- + Adaptation → FSM
- + Abstraction → Logical Tiles
- + Concurrency control → MVCC